

The 2026 prompting playbook — one page

From Prompts to Agents · researched Sep 2026 · full sources in the PLAYBOOK tab of your Course Workbook

TO DO — reliably helps

1	State the task precisely — verb, constraints, success criteria, and the WHY behind each rule	<i>Anthropic 2026 · Yang 2026: unstated requirements guessed right only 41.1%</i>
2	Structure with delimiters (XML/Markdown) and standardize ONE tested, versioned template	<i>Anthropic/OpenAI/Google guides · He 2024: wrapper alone swings up to 40%</i>
3	Zero-shot first; add 3-5 targeted, diverse examples only when format or tone matters	<i>OpenAI reasoning best practices · Anthropic 2026</i>
4	Long inputs: documents at the TOP, instructions at the END — bookend BOTH ends when very long	<i>Anthropic: up to ~30% · GPT-4.1 guide 2025: both ends beat either alone</i>
5	Give an out ("if it doesn't say, say so") + require citations	<i>Omar 2025: hallucination 66%→44% · Anthropic docs</i>
6	Concrete numeric budgets for measurable outputs — words, bullets, tool calls	<i>GPT-5.1 guide: adheres well to concrete length guidance · GPT-5 guide</i>
7	Persona + audience for VOICE and level — never for accuracy	<i>PersonaLLM 2024: traits detectable up to 80% · Google PTCF</i>
8	Self-check against NAMED criteria, with evaluation blinded from generation	<i>CoVe 2024: FactScore 55.9→71.4 · Cheng, Science 2026 (blinding)</i>

NOT TO DO — hurts or wastes effort

1	Contradictory instructions — reasoning models burn tokens reconciling them	<i>GPT-5 guide 2025</i>
2	Tips, threats, deadline pressure — null on average, ±35% per-question chaos	<i>Wharton R3 2025 · Salinas & Morstatter 2024</i>
3	Bare "do not hallucinate" commands — they trigger over-refusal of facts that ARE in the context (the "Safety Tax"); use the positive out+cite pattern instead	<i>arXiv:2601.02023 (Jan 2026)</i>
4	Revealing your preferred answer, or "are you sure?" as verification	<i>Cheng, Science 2026: +49% affirmation · Sharma 2023</i>
5	Piling micro-rules — even frontier models degrade past ~150 simultaneous instructions, early rules win	<i>IFScale, arXiv:2507.11538 · GPT-5.5 guide: smallest prompt that preserves the contract</i>
6	ALL-CAPS / ALWAYS / NEVER as an emphasis crutch — reserve absolutes for true invariants	<i>GPT-5.x guides · Gemini 3 guide</i>
7	Demanding JSON/schemas a human will just read — use structured output only when a system consumes it	<i>OpenAI 2024 · (do NOT teach "JSON hurts reasoning" as fact — 2025 replications)</i>
8	Concluding "the model can't do X" from one phrasing — formatting alone swings up to 76 points	<i>Sclar ICLR 2024 · Wharton R1</i>

EXPIRED — do the replacement instead

1	"Let's think step by step" / manual chain-of-thought → pick a reasoning model, set the effort dial	<i>Kojima 2022 origin (GSM8K 10.4→40.7 — it WAS real) · OpenAI now: avoid CoT prompts · Wharton R2</i>
2	Piles of few-shot exemplars (10+) → zero-shot first, then 3-5 format-definers	<i>Brown 2020 origin · DeepSeek-R1 paper 2025: few-shot degrades reasoners</i>
3	"You are a world-class expert" for accuracy → task-specific instructions; persona for voice only	<i>Zheng 2024 (162 personas: null) · Wharton R4 Dec 2025</i>
4	"Take a deep breath" and other magic phrases → automated prompt optimization (the method won; the phrase died)	<i>OPRO ICLR 2024</i>
5	Emotional appeals ("important to my career") → clarity; real stakes stated as factual context	<i>EmotionPrompt recalc arXiv:2409.20303: honest average ~2.6%, null on modern models</i>
6	"Instructions, ###, then text" + chopping documents for tiny windows → docs top, query end, bookend long context	<i>Liu 2023 (mechanism persists) · GPT-4.1 guide</i>
7	Hand-run self-consistency (10 samples, majority vote) + micro-decomposing every task → effort settings; chain only for auditable intermediates	<i>Wang ICLR 2023 origin · Anthropic 2026 · GPT-5.5: clear destination, let it choose the path</i>
8	Carrying your GPT-4-era prompt stack to each new model → re-baseline and re-test on every upgrade (PDCA)	<i>GPT-5.5 guide via Willison, Apr 2026</i>

Coherence: caps constrain quantity, absolutes lock judgment · repeat verbatim at both ends of long context or not at all · task detail helps, rule count hurts · CoT expired on frontier reasoning models only · "give an out" is "don't hallucinate" said positively.