

One skill, two modes: prompts that write — and prompts that work

GENERATIVE — a prompt that writes

you ask → it produces text → you read it and act

THE ANATOMY — SEVEN ELEMENTS (THE DESIGN LAYER)

Role — behaviours, not titles (“never invent numbers”). Shapes tone and framing, not IQ.

Task — one clear ask with a verb, the audience, and a success criterion.

Context — raw material, glossary, constraints; explain WHY a rule exists.

Format — structure, length cap, tone; say what TO do, not what to avoid.

Examples — 2–3 realistic samples of “good”, one edge case included.

The Out — “if the document doesn’t say, say so” — permission not to guess.

The Stop — where the action ends: “deliver X, then stop — nothing beyond.”

THE TOOLKIT IS A LOOP — PLAN · DO · CHECK · ACT

PLAN — design the ask: the seven elements, filled; be specific and say why; separate your ask from pasted material.

DO — run it in order: NUMBERED steps; “wait for my OK” where you want control; documents on top, ask at the end.

CHECK — verify before you trust: name the failed element; self-check against NAMED criteria (never “are you sure?”); blind review; reconcile numbers to a known total.

ACT — deliver + improve: fix ONE element and rerun; metaprompt; works twice? name it, version it; re-baseline on model upgrades.

Layers — you WRITE with the anatomy; you IMPROVE with the loop. In agentic work the loop is written INTO the brief — which is why it works best there.

EVIDENCE QUICK HITS

Works: specificity · examples · structure · grounding in documents · code for all math.

Myth: tips, politeness games, “genius” personas for accuracy — small, unstable effects.

Thinking models: skip “think step by step” and example piles; remove contradictions; set effort; stay clear.

Sycophancy: AI affirms you ~49% more than humans would. Never reveal your preferred answer; ask for the case AGAINST; paste your draft as “a colleague’s.”

TASK CHEATS — TEMPLATES G1-G8

Analysis (G2): exact columns + what one row means; all numbers via code; reconcile to a known total; numbered questions, then stop.

Writing (G1): feed raw material; hard word cap; “add no facts”; edit beats draft.

Evaluate (G3): anchored rubric; quotes as evidence; “not addressed = 1.”

Sheets (G4): name exact columns; README tab; formulas, not pasted values.

Decks (G5): titles = findings with numbers; approve text before rendering.

HTML (G6): single file, data embedded, works offline; show formulas on screen.

Research (G7): a question, not a topic; every bullet gets link + date; demand the gaps section.

Documents (G8): summaries serve a decision; verbatim quotes for anything contractual or numeric.

AGENTIC — a prompt that works

you brief a job → it plans, uses tools, checks itself, reports with evidence

THE MISSION BRIEF — 12 BLOCKS (TEMPLATE A1)

<role> behaviours, not job titles

<mission> goal, audience, deadline, “done looks like...”

<context> glossary, quirks, targets — the tribal knowledge

<environment> read-only inputs / outputs / forbidden actions

<inputs> each source: location, format, grain, trust level

<plan> steps in order; methods allowed vs ask-first

<checks> HARD = stop the run · SOFT = flag and continue

<outputs> exact filenames + CHANGES / FINDINGS / OPEN_ITEMS logs

<process> phases with human gates

<rules> never invent data; version outputs; log assumptions

<quality_bar> definition of done as an auditable checklist

<reporting> status · headline · evidence · open items, under 25 lines

Design thick, ship lean: design against all 12 — SHIP the four-section PLAN·DO·CHECK·ACT brief, one screen long.

GATES — CATCH ERRORS AT THE CHEAP END

Gate 1: definitions confirmed before any data is touched.

Gate 2: numbers reconciled to a total you already trust.

Gate 3: headlines approved as plain text before anything renders.

Autonomy: between gates, proceed. Stop and ask if: a HARD check fails · an input isn’t as described · a definition is ambiguous · a published number would change.

GUARDRAILS

Reversible → proceed. Irreversible (delete / send / publish / overwrite) → ask first.

Least privilege: the minimum folders, tools, and accounts for the job.

Anything the agent reads can carry hostile instructions (prompt injection). Highest risk when private data + untrusted content + outbound channels combine.

Vet third-party skills/plugins like software. Demand evidence of “done,” then spot-check.

A gate in the prompt is a suggestion; a gate in the harness is a control — put approvals in tool permissions too.

MANAGE PROMPTS LIKE ASSETS

Ladder: one-off → personal doc → team library → skills/instructions → repo (as code).

Trigger: explained the same task more than once? Package it.

Each entry: name · owner · version + change note · model tested · a filled example.

Eval lite: three gold cases (typical · edge · should-abstain) + a rubric; a new version must BEAT the current one; re-run on model upgrades.

Governance: prompts contain data — classify the library like the documents it quotes; secrets/PII never in a template ({{placeholders}} at run time).