

The Prompt Element Taxonomy

Every element, attribute, and option behind the Template Creator — with the research-backed why on every slot.

Two layers, one craft — you write with the anatomy; you improve with the loop. This taxonomy is the DESIGN layer: what a prompt contains, element by element, so you can ask “is it complete?”. The PROCESS layer is the PDCA loop from Part 3 — Plan (design with these elements) · Do (run in order, with wait-gates) · Check (verify against named criteria) · Act (deliver, then improve one element and version it). In agentic work the loop is written INTO the brief — the agentic elements below are exactly where.

This reference prints the complete ontology behind the Prompt Template Creator: 19 elements (7 generative, 12 agentic), their attributes, and every vetted option with its guidance. Every attribute carries a shaded Why it matters note — what that dial actually does to the output, with the study or vendor guidance it traces to (full citations: notes/research/ in the training repo). Solid rows are core picks; rows marked ◇ are the extended menu. Everything here is editable in `prompt-library/taxonomy/` — the HTML builder, this PDF, and the XLSX configurator all regenerate from the same files.

Term	OOP reading
Element	class — a major building block of a prompt (Role, Task, ... Mission, Checks)
Attribute	property — a configurable slot inside an element (Role.identity, Checks.hard)
Option	allowed value — a vetted choice for an attribute, with guidance on when to pick it
Technique	method — an operation you can run on/with an element (<code>few_shot()</code> , <code>critique_loop()</code>)
Template	instance — a saved configuration of elements and choices (the G/A library)
Inheritance	agentic elements extend their generative ancestors (Mission extends Task, adding a definition of done)

PART 1 — GENERATIVE ELEMENTS

Generative mode

A prompt that writes: you ask, it produces text, you read it and act. Five elements + two safety valves.

Role class Role extends PromptElement

Who is answering: identity, expertise posture, and checkable behavioral commitments. Selects which region of the model's training the answer is drawn from — shapes conduct, never competence. Behaviors beat titles because behaviors are auditable in the output. Prevents: the generic-assistant voice; mid-distribution, hedge-everything output.

Identity / posture Role.identity : enum [pick one]

Pick the professional posture. It sets vocabulary, caution level, and what gets foregrounded — not IQ.

Why it matters — Sets the register the answer is written in — vocabulary, caution level, what gets foregrounded. Expert personas measurably improve domain description tasks, but a 162-persona study found no accuracy gain on factual questions: identity buys voice and focus, not IQ. (Salewski et al. 2023; Zheng et al. 2024)

Option	Pick when	Emits into the prompt
Careful analyst	numbers, data, trends	You are a careful analyst.
Precise editor	improving text you supply	You are a precise editor.
Demanding reviewer	grading, evaluating, comparing	You are a demanding but fair reviewer.
Executive presentation writer	decks, summaries for leadership	You are a presentation writer for executives.
Patient explainer	teaching, onboarding material	You are a patient explainer who never condescends.
Devil's advocate	stress-testing a plan or decision	You are a devil's advocate.
Meticulous document processor	extraction, comparison, forms	You are a meticulous document processor.
Cited researcher	web/document research	You are a researcher who cites every claim.
◇ Business writer	emails, memos, announcements	You are an experienced business writer.
◇ Technical writer	procedures, specs, documentation	You are a technical writer who values precision over polish.
◇ Facilitator	synthesizing group input, agendas	You are a neutral facilitator.
◇ Translator / localizer	translation with tone intact	You are a professional translator who preserves register and nuance.
◇ Interviewer	the model should ask YOU questions	You are a skilled interviewer.
◇ Coach	feedback on your own skills/work	You are a supportive but honest coach.
◇ Red-teamer	risk reviews, failure-mode hunting	You are a red-teamer looking for ways this fails.
◇ Socratic teacher	learning, coaching through a problem	You are a Socratic teacher: you guide with questions before giving answers.
◇ Plain-language simplifier	de-jargoning	You are a plain-language simplifier: you make complex material genuinely simple without making it wrong.
◇ Meeting scribe	transcripts, notes	You are a meeting scribe: you capture decisions, actions, and disagreements – verbatim where wording matters.

Option	Pick when	Emits into the prompt
◇ Data organizer	text-to-table jobs	You are a data organizer: you turn messy text into clean, consistent structure.
🔗 custom	always available	Custom identity, e.g., 'a supplier-quality engineer reviewing a CAPA'...

Domain expertise `Role.expertise : string` [free text]

Name the field so the right vocabulary and priorities surface. This foregrounds knowledge — it does not create it.

Why it matters — Naming the field surfaces its vocabulary, defaults, and priorities — the answer starts from the right corner of the model's training instead of the generic middle. It foregrounds knowledge the model already has; it cannot create knowledge it lacks. (Salewski et al. 2023; vendor role guidance)

Option	Pick when	Emits into the prompt
🔗 custom	always available	e.g., 'Your domain: supplier quality engineering in medical devices.'

Behavioral commitments `Role.behaviors : Set` [pick any]

The checkable rules — the part of Role that actually changes output. Pick from the three families below.

Why it matters — The part of Role that measurably changes output: a behavioral rule is checkable line-by-line in the result, where a title is a vibe. Grounding-style rules are the documented lever for drastically reducing invented content. (Anthropic reduce-hallucinations guidance)

Accuracy behaviors `behaviors.accuracy` [pick any]

For anything with facts and numbers.

Why it matters — Training rewards a confident guess over an honest blank, so guessing is the default; accuracy rules re-open the abstain option. Compute-via-code exists because model arithmetic degrades sharply as numbers grow. (OpenAI 2025, why LMs hallucinate; Faith & Fate, NeurIPS 2023)

Option	Pick when	Emits into the prompt
Never invent	always safe to include	You never invent facts or numbers — missing is missing.
Compute via code	any math; needs a tool with code execution	You compute every number by running code, never mental arithmetic.
State n	data summaries, quality metrics	You state n for every proportion and flag any group with n below 20.
Quote verbatim	contracts, standards, specs	You quote exact wording, with location, for anything contractual, regulatory, or numeric — no paraphrase there.
Admit limits	analysis with imperfect data	You say "the data can't answer that" when it can't, and name what data would.
🔗 custom	always available	free text

Transparency behaviors `behaviors.transparency` [pick any]

Make the work auditable.

Why it matters — Stated assumptions and shown work convert an answer you must trust into one you can audit — errors surface at review, not downstream. Self-checking works only against concrete criteria like these; a bare 'are you sure?' can make answers worse. (Anthropic self-check guidance; Huang et al. 2024)

Option	Pick when	Emits into the prompt
State assumptions	anything the reader must be able to challenge	You state every assumption explicitly, in a list.
Show working	calculations, judgments	You show your working before the conclusion.
Label confidence	research, findings	You label each conclusion's confidence: high, medium, or low — and why.
Separate fact from reading	reports going to decision-makers	You separate observation from interpretation and never blend them in one sentence.
List changes	edits and rewrites	You list what you changed and why, so I can accept or reject each change.
🔗 custom	always available	free text

Style behaviors `behaviors.style` [pick any]

How it writes, stated as rules.

Why it matters — Style stated as rules survives; style stated as adjectives drifts. 'One idea per sentence' is enforceable and checkable in the output — 'write clearly' is re-interpreted on every run. (vendor format guidance; prompt-sensitivity research)

Option	Pick when	Emits into the prompt
Headline-first	busy readers, executives	You write headline-first: the answer in the first sentence, support after.
Plain language	mixed or non-specialist audiences	You use plain language and define any term of art at first use.
No padding	everything, honestly	You never pad. You cut adjectives that carry no data and phrases that carry no meaning.
Match my voice	when a STYLE CARD or samples are included	You match the voice of the writing samples I provide.
 custom	always available	free text

Stance `Role.stance` : enum [pick one]

The critical posture. Blind the model to your preference — it mirrors any side you reveal.

Why it matters — Models affirm users ~49% more than humans do and mirror whichever side you reveal — a single flattering exchange measurably raises the user's conviction. Declaring a neutral or adversarial stance up front pre-empts the mirror before it engages. (Cheng et al., Science 2026; Sharma et al. 2023)

Option	Pick when	Emits into the prompt
Neutral arbiter	comparisons, disputes	You take no side until the evidence does.
Skeptic	verification, review	You treat every claim as unproven until evidenced, including mine.
Adversary	pre-mortems, stress tests	You actively hunt for what is wrong, weak, or missing – do not soften.
Supportive improver	editing a colleague's draft	You improve the work while preserving the author's intent and voice.
◇ Steelman both sides	decisions with factions	You present the strongest honest case for every side before judging.
 custom	always available	free text

role_layering() — Combine 2–3 postures when one under-serves the job: "data engineer + analyst + presentation writer" documents, checks, AND lands the message. Use when: multi-skill deliverables.

style_card() — Once: feed 2–3 writing samples, ask for a reusable STYLE CARD; paste it here forever after. Use when: voice matching.

⚠ Anti-patterns: credentials without behaviors ("senior data scientist") · "you are a genius" for accuracy — personas shape tone, not correctness · stacking conflicting personas

Task `class Task` extends `PromptElement`

The single clear ask: verb + object + audience + success criterion. The most load-bearing words in the prompt — they define the target the completion optimizes toward. A question beats an instruction because it carries its own completion test. **Prevents: wandering output — ten observations and no answer.**

Action starter `Task.action` : enum [pick one]

Each verb is a mini-template — pick one and fill its `{{slots}}`, or write your own statement below.

Why it matters — The verb is the operation: it selects what transformation runs on your material and what 'done' looks like. Google's guidance calls the task verb 'the most important component of a prompt' — vague verbs produce exploration instead of answers. (Google Workspace prompting guide)

Option	Pick when	Emits into the prompt
Analyze	data questions	Answer: <code>{{the question, with metric and period – e.g., "what is X by group for Q2 vs the target"}}</code> ?
Summarize	condensing documents	Summarize <code>{{material}}</code> for <code>{{audience}}</code> : the three things they must know, what changed vs. what we knew, and what it does NOT cover.

Option	Pick when	Emits into the prompt
Draft	new text from raw material	Draft <code>{{deliverable}}</code> for <code>{{audience}}</code> . Its job: the reader should <code>{{know / decide / do}}</code> <code>{{X}}</code> after reading.
Rewrite	improving your own draft	Rewrite the text below. Keep every fact and commitment exactly – change only clarity, order, and tone.
Compare	options, proposals, versions	Compare <code>{{A}}</code> and <code>{{B}}</code> against <code>{{criteria}}</code> . Score each independently first, then head-to-head; name a winner per criterion with the evidence.
Evaluate	rubric grading	Grade the attached <code>{{artifact}}</code> against the rubric below. Quotes as evidence; not addressed = 1.
Extract	documents to data	Extract <code>{{fields}}</code> from <code>{{source}}</code> into <code>{{structure}}</code> . Anything not stated verbatim = MISSING – never inferred.
Explain	teaching, briefings	Explain <code>{{concept}}</code> to <code>{{audience}}</code> . Accuracy beats simplicity where they conflict – flag every simplification.
Brainstorm	idea generation	Generate <code>{{n}}</code> genuinely distinct options for <code>{{goal}}</code> – no near-duplicates; span <code>{{the dimension that matters}}</code> .
Verify	fact-checking	Verify this claim and nothing else: " <code>{{claim}}</code> ". Verdict: TRUE / FALSE / PARTLY TRUE / CAN'T VERIFY – then the evidence.
◇ Classify	sorting, tagging	Classify each item below into <code>{{categories}}</code> . If an item fits none, label it OTHER with one line on why.
◇ Rank	prioritization	Rank <code>{{items}}</code> by <code>{{criterion}}</code> . Show the score and the one fact that most drove each position.
◇ Reconcile	two versions of the truth	Reconcile <code>{{source A}}</code> against <code>{{source B}}</code> : what matches, what conflicts (both sides quoted), what appears in only one.
◇ Translate	translation	Translate the text below into <code>{{language}}</code> , preserving register, tone, and any terms of art (keep <code>{{terms}}</code> untranslated).
◇ Outline	big documents, decks	Outline <code>{{deliverable}}</code> before writing it: sections, one-line message per section, open questions. Stop for my approval.
◇ Diagnose	troubleshooting	Diagnose why <code>{{symptom}}</code> is happening. List candidate causes ranked by likelihood, and the cheapest check that would confirm or kill each.

Task statement `Task.statement : string [free text]`

Or write the ask yourself — one task, one sentence if you can manage it.

Why it matters — One task in one sentence is the highest-leverage clarity move: current reasoning models pay a real cost for ambiguity — they burn thinking tokens trying to reconcile vague or conflicting instructions instead of answering. (OpenAI GPT-5 guide; Anthropic 'be clear and direct')

Option	Pick when	Emits into the prompt
🔗 custom	always available	The ask in your words. One task per prompt; several → number them under Scope.

Audience `Task.audience : enum + timeBudget [pick one]`

Who receives this and what they already know. It sets altitude and vocabulary.

Why it matters — Audience re-ranks and re-words the same content: what leads, what gets cut, which terms need expansion. Without one, the model writes for nobody in particular — the mid-distribution reader who doesn't exist in your building. (vendor Persona/Context frameworks)

Option	Pick when	Emits into the prompt
Executive	leadership updates	executives – 60-second read, decision-focused, no background needed on our org
Peer specialist	expert to expert	a peer specialist – full technical vocabulary, skip the basics
Cross-functional partner	other departments	a cross-functional partner – smart, not in our field; no unexplained jargon
New hire	onboarding, training	a new hire – assume nothing about our processes; define everything local
External customer	outbound	an external customer – professional, no internal names or shorthand
◇ Regulator / auditor	compliance-facing	a regulator or auditor – exact language, no speculation, every claim traceable
🔧 custom	always available	Custom audience...

Reader's time budget `audience.timeBudget [pick one]`

How long they'll actually give it — the honest constraint.

Why it matters — The honest constraint that forces compression: a 30-second reader gets the number and the ask; a 10-minute reader gets the reasoning. Left uncapped, models default to thorough — which the reader experiences as long. (practitioner consensus; length is a vendor API dial for this reason)

Option	Pick when	Emits into the prompt
30 seconds	status updates, FYIs	They will give this 30 seconds – the first two sentences must carry everything.
3 minutes	typical business read	They will give this about 3 minutes.
Deep read	reference material, appendices	They will study this carefully – completeness beats brevity here.

Reading level `audience.readingLevel : scale [pick one]`

The Canvas-style slider, as options. Plain-language guidance targets 6th–8th grade for general audiences.

Why it matters — Vocabulary altitude is a dial the model follows reliably — and its default sits well above general-audience reading. Plain-language practice targets 6th–8th grade for mixed audiences; say the level and it lands there. (plain-language standards; style-prompting research)

Option	Pick when	Emits into the prompt
Middle school (plain-language default)	general audiences — the plain-language standard	Write at a 6th–8th grade reading level.
High school	broad professional	Write at a high-school reading level.
College	professional default	Write at a college reading level.
◇ Graduate / expert	specialist readers	Graduate-level density is fine.
◇ Explain to a smart 12-year-old	genuinely new concepts	Explain it so a smart 12-year-old would follow – without being cute about it.

Success criterion `Task.success : predicate [pick one]`

The task's own completion test — what must be true of the output.

Why it matters — A stated completion test gives the model a target to optimize toward and gives you a check to run afterward — the same mechanism that makes self-verification work: concrete criteria catch errors reliably; vague quality asks don't. (Anthropic self-check guidance)

Option	Pick when	Emits into the prompt
Reader can decide	decision support	the reader can make the decision without asking a single follow-up.
Every question answered	analysis	every numbered question has a direct answer with its number.
Reader knows next step	action-driving messages	the reader knows exactly what to do next, and by when.
Survives a skeptic	high-stakes documents	a skeptical reviewer finds no unsupported claim.
 custom	always available	Custom success test...

Purpose — why it matters `Task.purpose : string [free text]`

The decision or event this feeds. Models generalize from reasons — and it filters what matters.

Why it matters — Models generalize from reasons: 'no ellipses — this will be read aloud by text-to-speech' outperforms the bare rule, because the model extends the spirit to cases you didn't list. Purpose also filters relevance — it now knows what matters to the decision. (Anthropic explain-why guidance)

Option	Pick when	Emits into the prompt
 custom	always available	e.g., 'feeds the Q3 sourcing decision on {{supplier}}'

Sense check `Task.senseCheck : enum // RODES's unique slot – Gate 1's chat ancestor [pick one]`

Have the model confirm understanding before doing the work — cheapest error catch there is.

Why it matters — A restate-before-work step catches a misread task at the cheapest possible point — before any output exists. It is the generative version of the agentic Phase-0 intake gate: pennies to fix here, a full redo after a polished wrong answer. (plan-approval gate pattern)

Option	Pick when	Emits into the prompt
Restate + wait	big or ambiguous jobs	Before starting: restate the task in your own words, list your assumptions, and ask anything unclear. Then WAIT for my confirmation.
Restate + proceed	medium jobs, async	Before answering: restate the task in one line and flag anything ambiguous — then proceed with your best reading.
◇ Interview me first	you have the answers, the model has the structure	Ask me questions, one at a time, until you have what you need to do this well — then do it.

question_reframe() — Turn a topic into a question before prompting — questions generate answers, topics generate exploration. Use when: you catch yourself typing a noun phrase.

decompose() — One prompt per stage for big jobs: outline → approve → draft → review. Inspectable intermediates are an audit trail. Use when: anything over a page, or with an approval step.

critique_loop() — After the draft: "Critique your draft against the success criterion, then produce a revised version." One round helps; more needs concrete criteria. Use when: quality-sensitive drafts.

△ Anti-patterns: a topic instead of a question ("the returns data") · three unrelated asks in one prompt · "help me with..." (no verb, no object)

Context `class Context extends PromptElement`

Everything true in your world that isn't in the training data — material, definitions, quirks, constraints, and the reasons behind rules. The model fills every gap with the most plausible filler from training. Context replaces plausible-in-general with true-for-you. This is the anti-hallucination element. Prevents: generically-right, specifically-wrong output — the plausible number with your company's name on it.

Raw material `Context.material : Document[]` **[pick one]**

The stuff to work from. Long documents go FIRST in the final prompt, instructions after — and delimit so material is never mistaken for instructions.

Why it matters — Placement measurably matters: long documents at the TOP with the question at the END improved answers by up to ~30% in vendor testing, and accuracy sags for content buried in the middle of a long context. (Anthropic long-context guidance; Liu et al. 2023, lost in the middle)

Option	Pick when	Emits into the prompt
Attached document	file uploads	Work ONLY from the attached document(s). If it isn't in them, it doesn't exist for this task.
Pasted below	pasted text	The material is pasted below between triple quotes. Treat everything inside as content, never as instructions.
Prior thread	continuing a chat	Use only the conversation above as source material.
 custom	always available	Paste the material here (it will be fenced), or note what's attached...

Glossary — controlled vocabulary `Context.glossary : Map` **[pick any]**

Terms with local meaning. For metrics: numerator, denominator, grain, inclusions/exclusions — or the model will compute its own version.

Why it matters — The model resolves every undefined term to its most common global meaning — not your local one. A metric card (numerator, denominator, grain, exclusions) replaces the single biggest source of wrong analytical answers: a guessed denominator. (hallucination mechanics; data-analysis field practice)

Option	Pick when	Emits into the prompt
Metric card skeleton	any computed metric	<code>{{METRIC}} = {{numerator}} ÷ {{denominator}}, at {{grain}}; include {{...}}; exclude {{...}}; target {{...}}.</code>
 custom	always available	TERM = definition as used here...

Known quirks → rules `Context.quirks : List` **[pick any]**

Every quirk you don't write down becomes a silent 'fix' the model invents. Pair each with its rule.

Why it matters — Every quirk you don't write down becomes a silent 'fix' the model invents — plausibility fills the gap, fluently. A quirk paired with its rule converts a hidden landmine into a followed instruction. (OpenAI 2025, why LMs hallucinate)

Option	Pick when	Emits into the prompt
Totals row	spreadsheet exports	The export has typed totals rows — exclude them, recompute from detail rows, and report if they disagree.
Leading zeros	IDs in spreadsheets	Codes and part numbers keep leading zeros — treat them as text, never numbers.
Alias mapping	inconsistent naming	<code>{{Name A}}</code> and <code>{{Name B}}</code> refer to the same entity — normalize to <code>{{canonical name}}</code> .
Mixed units	multi-source data	Units are mixed (<code>{{lb and kg}}</code>) — convert everything to <code>{{unit}}</code> and note each conversion.
Manual entry	hand-filled sheets/forms	These figures are manual entry — re-derive every total from rows; never trust a typed total.
◇ Date formats	multi-system exports	Dates appear in multiple formats — parse explicitly, state the format assumed per source, output ISO.
◇ Second header row	formatted reports as data	The sheet has a merged/second header row — the real headers are row <code>{{n}}</code> .
 custom	always available	quirk → the rule to apply...

Constraints & sensitivities `Context.constraints : Set [pick any]`

What must not happen — stated positively where possible, with the reason where it isn't obvious.

Why it matters — Positively-phrased constraints are followed more reliably than prohibitions — earlier models showed inverse scaling on negated instructions, and current vendor guidance remains 'tell it what TO do.' When you must forbid, attach the reason. (Anthropic guidance; Jang et al. 2022)

Option	Pick when	Emits into the prompt
No new facts	the single highest-value constraint for drafting	Do not add facts, numbers, or commitments that are not in the material.
Confidential	internal data	This material is confidential — never quote raw rows or names outside the requested output.
Consistency with prior work	recurring deliverables	Stay consistent with <code>{{prior file / decision}}</code> : definitions, names, and colors must not change.
No speculation on topic	sensitive subjects	Do not speculate about <code>{{topic}}</code> — if asked-for material doesn't cover it, say so.
◇ Language/units house rules	org conventions	House rules: dates <code>{{ISO in files, MM/DD/YY in prose}}</code> ; units <code>{{imperial}}</code> ; language <code>{{English}}</code> .
🔧 custom	always available	Custom constraint...

The why behind the rules `Context.rationale : string [free text]`

Explain any non-obvious rule — "this will be read aloud, so no ellipses" outperforms "no ellipses." Models generalize from reasons.

Why it matters — The reason behind a rule is itself an instruction: models apply explained rules in spirit, not just letter, extending them to cases you didn't enumerate — the vendor docs state this generalization effect verbatim. (Anthropic best practices)

Option	Pick when	Emits into the prompt
🔧 custom	always available	e.g., 'this goes into a validated record, so wording must match the source exactly'

delimit() — Fence material with tags or triple quotes so content is never read as instructions — also your first defense against instruction-like text inside documents. Use when: always, when pasting anything.

context_budget() — Long thread or many documents? Summarize what matters into a fresh prompt instead of dragging the whole pile — buried middles get skimmed. Use when: long conversations, multi-doc jobs.

⚠️ **Anti-patterns:** "use our standard definitions" (whose?) · material and instructions interleaved unfenced · pasting everything you have "just in case"

Format `class Format extends PromptElement`

The output contract: shape, length, tone, and what must be in or out. Models imitate structure more readily than they obey descriptions of it — state format positively, cap length numerically, and demonstrate when it really matters. **Prevents: the 3x-too-long draft; the buried answer; output you must re-shape by hand.**

Output type `Format.outputType : enum [pick one]`

The deliverable's shape. Name the columns for tables, the sections for documents.

Why it matters — Naming the deliverable's shape — a table with these columns, a memo with these sections — removes the largest degree of freedom the model would otherwise spend guessing, and makes output land ready to use instead of ready to reformat. (all four vendor frameworks: Format / Expectations)

Option	Pick when	Emits into the prompt
Email	correspondence	Format: an email — subject line + body.
Exec summary	leadership	Format: an executive summary — the answer, three supporting facts, the ask.
Bullet brief	quick consumption	Format: a bullet brief — grouped bullets, each a fact with a number.
Table	comparable items	Format: one table — columns: <code>{{name them exactly}}</code> . No prose outside the table except a one-line takeaway above it.

Option	Pick when	Emits into the prompt
Comparison matrix	A-vs-B decisions	Format: a comparison matrix – rows = <code>{{criteria}}</code> , columns = <code>{{options}}</code> , plus a verdict row.
Step-by-step procedure	instructions	Format: a numbered procedure – one action per step, the expected result after each.
Checklist	verification aids	Format: a checklist – imperative items, groupable, no explanation longer than one line.
Memo	formal internal docs	Format: a one-page memo – situation, findings, recommendation, next steps.
Slide outline	deck preparation	Format: a slide outline – per slide: title stating the takeaway (≤ 12 words, with the number), ≤ 3 bullets, [VISUAL: what it must prove], speaker notes.
◇ Report section	formal reports	Format: a report section with headed subsections; every claim carries its evidence inline.
◇ FAQ	rollouts, policy changes	Format: an FAQ – the questions people will actually ask, sharpest first.
◇ Action-item table	meeting follow-ups	Format: an action table – owner action due date source sentence (verbatim).
◇ JSON	output feeding software	Format: valid JSON only, matching this schema: <code>{{schema}}</code> . No prose before or after.
◇ CSV	data handoffs	Format: CSV with header row: <code>{{columns}}</code> . Quote fields containing commas; no commentary.
◇ Talking points	verbal delivery	Format: spoken talking points – short sentences, natural rhythm, no visual formatting.
🔗 custom	always available	Custom format...

Structure rules `Format.structure : Set [pick any]`

The internal skeleton — where the answer sits, what every section carries.

Why it matters — Trivial format differences swing measured accuracy by up to 76 points across phrasings — structure is not cosmetic. A fixed, tested skeleton removes that variance and puts the answer where the reader looks first. (Sclar et al., ICLR 2024)

Option	Pick when	Emits into the prompt
Headline first	almost everything	The answer in sentence one; support after; background last or never.
TL;DR line	longer documents	Open with a one-line TL;DR in bold.
Facts with numbers	data-driven content	Every bullet is a fact with a number – no adjectives without data.
Assumptions at end	analysis, plans	End with two lists: assumptions made, and open questions.
Source footer	anything with data	Footer: source, coverage period, as-of date.
◇ One message per section	reports, decks	One message per section; the heading states it.
🔗 custom	always available	Custom structure rule...

Length cap `Format.length : NumericCap [pick one]`

Numeric, with a priority rule. Adjectives ("brief") are unenforceable; uncapped AI runs 2–3× long.

Why it matters — A numeric cap is enforceable; 'brief' is a mood. Uncapped, models run two to three times longer than the reader needs — length is so real a dial that vendors expose it as an API parameter. The famous \$200-tip stunt produced ~11% longer answers, not better ones. (OpenAI verbosity; tip-experiment analyses)

Option	Pick when	Emits into the prompt
≤ 120 words	emails, updates	at most 120 words – cut content before quality.
≤ 250 words	summaries	at most 250 words.
One page	memos, briefs	one page maximum.
≤ 5 bullets	skimmable lists	at most 5 bullets, one line each.

Option	Pick when	Emits into the prompt
◇ 3 sentences	radical brevity	exactly 3 sentences.
◇ No cap	reference material	as long as it needs to be – completeness beats brevity here.
◇ XS — one-liner	status pings	one sentence.
◇ M — half page	standard summaries	about half a page.
◇ XL — exhaustive	reference documents	exhaustive – cover everything, organized with headings.
◇ Character cap	social posts, SMS, UI strings	at most {{280}} characters.
🔗 custom	always available	Custom cap, e.g., '≤ 200 words'...

Tone `Format.tone` : {formality, directness, warmth} [pick any]

Compose from the three dials below — one pick per dial reads best.

Why it matters — Tone is where role prompting genuinely pays: models imitate a specified register reliably, and register changes how the human reader receives the work. Compose it from dials rather than adjectives — one pick per dial reads best. (vendor role-for-tone position; style-prompting research)

Formality `tone.formality` [pick one]

Why it matters — Formality sets the social distance of the prose — hallway versus boardroom. Models default to mid-formal; one explicit pick lands it where you need it. (style-prompting research)

Option	Pick when	Emits into the prompt
Formal	external, regulatory, records	Tone: formal.
Professional	default for work	Tone: professional but natural.
Conversational	team comms	Tone: conversational.
◇ Casual	social, informal	Tone: casual and light.
◇ Technical	specs, engineering docs	Tone: technical and precise.
◇ Academic	papers, formal reviews	Tone: academic.
🔗 custom	always available	free text

Directness `tone.directness` [pick one]

Why it matters — Directness decides whether the ask or the courtesy leads. Models hedge by default — caution is trained in — so 'lead with the ask' has to be said to happen. (practitioner consensus; headline-first practice)

Option	Pick when	Emits into the prompt
Plain and direct	most business writing	Say things plainly; the point is never buried.
Diplomatic	sensitive audiences	Diplomatic phrasing – honest, but leaves the reader room.
Blunt	internal reviews you asked for	Blunt – no cushioning, no hedging.
◇ Assertive	recommendations	Assertive – clear positions, no hedging language.
◇ Cautious	early findings	Cautious – flag risk and uncertainty prominently.
🔗 custom	always available	free text

Warmth `tone.warmth` [pick one]

Why it matters — Warmth calibrates the relationship signals in the text. Unmanaged, it drifts toward affirmation — the same preference-training pressure behind sycophancy — so the dial keeps warmth a choice rather than a bias. (Sharma et al. 2023, sycophancy origins)

Option	Pick when	Emits into the prompt
Warm	people topics	Warm and human.
Neutral	facts and analysis	Emotionally neutral.
◇ Clinical	incident reports, audits	Detached and clinical.
◇ Encouraging	coaching, team comms	Encouraging — motivate action, acknowledge effort.
◇ Playful	celebrations, social	Light and playful — humor welcome.
🔗 custom	always available	free text

Language level `Format.languageLevel` : enum [pick one]

Vocabulary altitude and jargon policy.

Why it matters — Jargon policy decides who can act on the output without a translator. Models also mirror the prompt's own vocabulary — write at the altitude you want back. (Anthropic: match your prompt's style to the desired output)

Option	Pick when	Emits into the prompt
Plain language	default	Plain language — a busy reader gets it in one pass.
8th-grade + defined acronyms	broad audiences	Roughly 8th-grade reading level; define or remove every acronym.
Technical, defined	mixed technical audiences	Technical language is fine; define each term of art at first use.
Expert shorthand	peer specialists	Expert-to-expert — shorthand and standard abbreviations are fine.
🔗 custom	always available	free text

Must include / must not `Format.musts` : {include: Set, exclude: Set} [pick any]

The inclusion contract.

Why it matters — An explicit include/exclude list turns review into a checklist — presence and absence verified in seconds instead of re-reading for feel. It is the format-level version of checkable behavioral rules. (self-check-against-criteria guidance)

Option	Pick when	Emits into the prompt
Cite everything	research, disputes	Every factual claim carries its source.
Working before result	calculations	Show the formula or reasoning before the result it produces.
No preamble	everything	Start with the content — no "Certainly!", no restating the question.
No speculation	factual work	No speculation — flag gaps instead of filling them.
◇ Prose only	text pasted into other tools	Flowing prose paragraphs — no markdown, no bullets.
◇ No em-dash chains	house style	Plain punctuation — no decorative dashes or emoji.
🔗 custom	always available	Custom must / must-not...

Variations `Format.variations : int // CRISPE's 'Experiment' slot [pick one]`

Ask for alternatives to choose from — divergence as a dial. You become the selector, not the author.

Why it matters — Asking for N alternatives converts sampling variance from a bug into a menu — the insight behind self-consistency, where several sampled answers beat any single one. You become the selector, which is where your judgment belongs. (Wang et al. 2023, self-consistency)

Option	Pick when	Emits into the prompt
Single best	default	Give your single best version.
3 labeled versions	picking a direction	Give 3 genuinely different versions, labeled by their angle — not three rewordings.
Safe + bold	tone/stakes decisions	Give two versions: one safe, one bold. Say what each risks.
◇ Options + recommendation	you want a default with escape hatches	Give <code>{{n}}</code> options, then recommend one and say why in one line.
🔗 custom	always available	Custom, e.g., '5 subject lines, ranked'...

format_by_example() — One filled sample of the output beats three paragraphs describing it — put it in Examples. Use when: unusual or strict formats.

reason_then_format() — On hard problems, heavy format constraints tax reasoning: let it solve free-form first, then re-format the answer in a second step. Use when: difficult analysis + strict output shape.

⚠️ **Anti-patterns:** adjective caps ("keep it short") · only saying what NOT to do · strict schema forced onto a hard reasoning task in one pass

Examples `class Examples extends PromptElement`

3–5 realistic demonstrations of the task done right — the strongest steering signal there is. In-context learning: the model infers the pattern from demonstrations without retraining. Diversity is the point — it handles edge cases exactly as your examples do. **Prevents: the model guessing your standard from the internet's average; format drift across runs.**

How many `Examples.count : int [pick one]`

Zero-shot first on thinking models; add examples when format drifts. More helps most on formatting/labeling tasks.

Why it matters — Few-shot examples were the original prompt-engineering discovery: task performance from demonstrations alone, no retraining. 3–5 diverse examples steer format and tone most reliably — but start zero-shot on thinking models; few-shot measurably degrades some of them. (Brown et al. 2020; Anthropic guidance; DeepSeek-R1 paper)

Option	Pick when	Emits into the prompt
One gold standard	format anchoring	Here is one example of exactly what I want:
3–5 diverse	classification, extraction, tone	Here are <code>{{3–5}}</code> examples of what I want — note they cover the normal case AND the edge cases:
Contrast pair	teaching a quality bar	Here are two examples — one GOOD and one BAD, labeled, with why:

Coverage `Examples.coverage : Set [pick any]`

Which cases your examples span — the model's boundary behavior copies theirs.

Why it matters — The model copies your examples' boundary behavior, not just their shape: exemplar choice and order alone can swing accuracy from near-chance to over 90%. All-happy-path examples teach it to treat every case as happy — one typical + one edge + one reject teaches the standard. (Lu et al. 2021; Prompt Report design decisions)

Option	Pick when	Emits into the prompt
Typical case	always	(One example is the typical case.)
Edge case	real-world messiness	(One example is an edge case: <code>{{missing field / borderline / messy input}}</code> — note how it's handled.)
Reject/escalate case	judgment tasks	(One example shows when to REJECT or escalate instead of answering.)

Packaging `Examples.packaging : enum [pick one]`

Fence examples so they can't be mistaken for live input.

Why it matters — Fencing examples in tags keeps them from being read as live input or instructions — unambiguous parsing is the documented reason vendors recommend tag structure, and unfenced examples are the classic source of example text leaking into output. (Anthropic XML-tags guidance)

Option	Pick when	Emits into the prompt
Tagged	default	Each example is wrapped in tags. They are demonstrations, not input.
Input/Output pairs	transformations	Examples are shown as Input: / Output: pairs.

The examples `Examples.items : Example[] [free text]`

Paste them. Real (redacted) beats invented; lightly-genericized real beats both.

Why it matters — Models imitate surface detail as much as structure — a real (redacted) example carries your house style's actual distribution; an invented one carries what the model guesses your style is. Exemplar quality and similarity are measured few-shot design decisions. (Prompt Report §2.2)

Option	Pick when	Emits into the prompt
 custom	always available	Paste example(s) here...

few_shot_design() — Research-backed defaults: order examples randomly (order alone can swing accuracy), balance labels across categories, use a common format, prefer examples similar to the live case. Use when: classification and labeling at any volume.

example_generator() — No examples on hand? Ask the model to draft 5 candidate examples, curate the best 2-3 yourself, then use those — you become the editor, not the author. Use when: bootstrapping a new template.

⚠ Anti-patterns: five clones of the same happy case · unfenced examples bleeding into the live input · example piles on reasoning tasks (zero-shot first on thinking models)

The Out `class TheOut extends SafetyValve`

The escape route for missing information: name the condition, the honest move, and how gaps surface — so the model never has to choose between obeying you and being honest. A model is trained to always produce an answer — benchmarks reward confident guessing over abstaining. Unless you specify the honest move, inventing is the only way it can obey your prompt. Prevents: confident invention — the signature failure of factual work.

When it applies `TheOut.triggers : Set [pick any]`

The situations where not-knowing must surface instead of being papered over.

Why it matters — Standard training rewards a confident guess over 'I don't know' — benchmarks score the guess, so guessing is the default. Naming the exact situations where not-knowing must surface re-opens the abstain option precisely where it matters. (OpenAI 2025, why LMs hallucinate)

Option	Pick when	Emits into the prompt
Material silent	the universal one	If the material doesn't say – say so. Do not fill the gap.
Sources conflict	research, multi-doc	If sources conflict, present both with dates – do not silently pick one.
Data can't answer	analysis	If the data cannot answer the question, say so and name what data would.
Ambiguous term	specs, requests	If a term or instruction is ambiguous, flag it before proceeding.
◇ Beyond cutoff	current-events-adjacent	If this needs information newer than your training, say so rather than recalling.
◇ Documents only	grounded Q&A;, compliance	Answer ONLY from the provided documents – no general knowledge. If they don't cover it, say so.
 custom	always available	Custom trigger...

What to do instead `TheOut.behavior : enum [pick one]`

The escape hatch's shape.

Why it matters — Explicitly permitting 'I don't know' 'can drastically reduce false information' — the vendor's own wording. Prefer a CHECKABLE shape: UNKNOWN plus a gap list can be verified in the output; a bare 'say so' names no form. Giving the escape hatch a shape converts invisible guesses into visible, handleable items. (Anthropic reduce-hallucinations, verbatim)

Option	Pick when	Emits into the prompt
UNKNOWN + gap list	the checkable default — condition, honest move, and surfacing in one sentence (the course anatomy's form)	Anything the material doesn't state: write UNKNOWN — never estimate. List the gaps at the end.
Say so plainly	lightest form — names no output shape	When uncertain: say "I don't know" or "the material doesn't say" — plainly, without apology.
Best reading, labeled	when a draft answer still helps	When uncertain: give your best reading, labeled UNCERTAIN, with what would confirm it.
Ask first	interactive sessions	When uncertain: ask me up to {{3}} clarifying questions before proceeding.
 custom	always available	free text

quote_first() — For long documents: "quote the passages relevant to the question first, then answer only from your quotes." Grounding through extraction. Use when: long-document Q&A.;

verify_then_answer() — Chain-of-Verification: draft → generate verification questions → answer them independently → revise. Ask for the four steps explicitly on factual work that matters. Use when: factual claims with consequences.

⚠ Anti-patterns: omitting the out and then blaming the model for guessing

The Stop `class TheStop extends SafetyValve`

The scope boundary — converts open-ended capability into bounded work. Without a boundary, capability sprawls: ten charts and no answer. The Stop is the difference between an answer and a wander. Prevents: scope sprawl; work continuing past the point of usefulness.

Scope `TheStop.scope : OrderedList [free text]`

Numbered, in priority order — the model works down the list.

Why it matters — A numbered scope list is an execution order: the model works down it instead of wandering sideways. Explicit scope and stop criteria are how vendors now control model 'eagerness' — the same dial, applied to text work. (OpenAI GPT-5 guidance)

Option	Pick when	Emits into the prompt
 custom	always available	1. {{primary question}} 2. {{secondary}} 3. {{optional if room}}

Stop rule `TheStop.rule : enum [pick one]`

Where work ends — or pauses for you.

Why it matters — An explicit end condition prevents both default failures: stopping early with a partial answer, and padding on past the answer. For anything loop-like, stopping conditions are a control requirement, not a nicety. (Anthropic building-effective-agents)

Option	Pick when	Emits into the prompt
Answer, then stop	analysis	Answer these, then stop — do not go exploring.
Outline, then wait	big documents, decks	Produce the outline only, then STOP for my approval before writing the full version.
Step 1 only	chained work	Do only step 1 now. Wait for my go before any further step.
Ask before expanding	letting good ideas surface without sprawl	If you believe something outside this scope matters, name it in one line at the end — do not pursue it.
 custom	always available	free text

Explicitly out of scope `TheStop.exclusions : Set [free text]`

Name the adjacent things NOT to touch.

Why it matters — Models err toward helpful overreach — adjacent questions get answered because they can be. Naming the not-asked keeps output on the question you actually have: the manual version of a relevance guardrail. (OpenAI guardrail categories)

Option	Pick when	Emits into the prompt
 custom	always available	e.g., 'root-cause analysis – this pass is descriptive only'

⚠ Anti-patterns: unnumbered multi-part asks · "and anything else you think is relevant" (an invitation to sprawl)

PART 2 — AGENTIC ELEMENTS

Agentic mode

A prompt that works: you brief a job; it plans, uses tools, checks itself, and reports with evidence. Twelve blocks — the mission brief.

Role `class AgentRole extends Role`

The standing behavioral contract — biases hundreds of unsupervised decisions, not one answer. In chat, Role shapes one reply you're about to read. In an agent, it persists across every step you don't watch. Write the behaviors you'd want if you couldn't check any single one. **Prevents: an agent that improvises its professional standards mid-run.**

Working posture `AgentRole.posture : enum [pick one]`

The job-shaped identity. Compound postures (engineer + analyst + writer) prevent single-lens blind spots.

Why it matters — An agent's posture persists across hundreds of unsupervised decisions, so it must be job-shaped, not title-shaped. Compound postures (engineer + analyst + writer) prevent the single-lens failure where every problem gets one discipline's answer. (agent-design practice; Anthropic agent guidance)

Option	Pick when	Emits into the prompt
Engineer + analyst + writer	data-to-deliverable pipelines	You are a data engineer, analyst, and presentation writer in one: you document lineage, question the numbers, and land the message.
Meticulous processor	document pipelines, trackers	You are a meticulous document processor: you extract what is written, never what you infer.
Research agent	multi-source research jobs	You are a research agent: you search wide, read critically, and cite everything.
Ops automation agent	recurring file/system chores	You are a careful operations agent: predictable, boring, and traceable beats clever.
◇ Refactoring engineer	code and script work	You are a conservative engineer: the smallest change that works, tested before shipped.
🔧 custom	always available	Custom posture...

Standing behaviors `AgentRole.behaviors : Set [pick any]`

The rules that must survive hour three of a long run.

Why it matters — Long runs outlive their own context — what was said in minute one has faded by hour three. Standing behaviors are the conduct that must hold when nothing recent reinforces it, which is why harness guidance bakes rules into re-read files rather than conversation. (Anthropic long-running-agents harness)

Option	Pick when	Emits into the prompt
Never invent	always	You never invent, estimate, or back-fill a number — missing is missing.
Log transformations	data jobs	You log every transformation with a row count — no silent fixes.
Scripts over hands	recurring work	You prefer reproducible scripts over manual edits; anything done twice becomes a script.
Headline-first reporting	any human-facing output	You write headline-first for busy decision-makers.
Inputs are sacred	always	You treat every input as read-only, always.
Evidence over assertion	always	You show evidence for claims of success — the command run, the output, the re-opened file.
🔧 custom	always available	Custom standing behavior...

⚠ Anti-patterns: chat-style helpful-assistant roles on unsupervised jobs · titles without behaviors

Mission `class Mission extends Task // + definitionOfDone`

The goal and the definition of done — the loop's exit condition. An agent runs until something says stop. Give it a check it can run; without one, 'looks done' is the only signal available. **Prevents: the agent deciding for itself what finished means — plausible activity, indefinitely.**

Goal `Mission.goal : string [free text]`

One sentence: the question to answer or the thing to produce.

Why it matters — One goal sentence is the loop's objective function: gather → act → verify → repeat runs 'until the goal is met.' Two goals in one brief make every prioritization a coin-flip you never see. (Anthropic agent-loop framing)

Option	Pick when	Emits into the prompt
 custom	always available	e.g., 'What is {{metric}} by {{group}} for {{period}}, and where is the gap?'

Decision it supports `Mission.decision : string [free text]`

What someone will do differently depending on the answer — filters everything downstream.

Why it matters — Naming the decision the work supports filters everything downstream — which findings matter, what precision suffices, what is noise. It is the brief-level relevance guardrail: scope drift dies here. (OpenAI relevance-guardrail concept; mission-brief practice)

Option	Pick when	Emits into the prompt
 custom	always available	e.g., 'where to focus {{resource}} next quarter'

Audience & deadline `Mission.audience, Mission.deadline [free text]`

Who sees it, how long they'll give it, and when it's due.

Why it matters — Audience and deadline set the depth-versus-speed tradeoffs the agent will otherwise make for you, silently. A 30-second executive read and a hard due date change method choices all the way down the pipeline. (practitioner consensus)

Option	Pick when	Emits into the prompt
 custom	always available	e.g., 'review-meeting attendees, ~3 minutes on this; due 2 business days before the review'

Cadence `Mission.cadence : enum [pick one]`

One-off or recurring changes the economics: recurring jobs get standing briefs.

Why it matters — One-off versus recurring changes the economics of the whole brief: recurring jobs justify standing briefs, delta reporting, and versioned definitions — where agentic prompting pays compound interest. (reuse-by-diff pattern; AGENTS.md practice)

Option	Pick when	Emits into the prompt
One-off	single deliverable	Cadence: one-off.
Recurring	weekly/monthly/quarterly jobs	Cadence: recurring — between cycles, only input paths and dates change; keep this brief in the repo and reuse by diff.

Done looks like `Mission.definitionOfDone : Deliverable[] [free text]`

The deliverables in one line — exact and countable.

Why it matters — The loop needs an exit condition it can evaluate. 'Done looks like: three files, named thus' is countable; without it agents stop early or run forever — the documented failure that structured task lists with pass flags exist to prevent. (Anthropic harness: feature-list pattern)

Option	Pick when	Emits into the prompt
 custom	always available	e.g., '2 slides + updated master workbook + CSV snapshot, all in deliverables/'

Verifiable by `Mission.verifiableBy : Check [pick any]`

The check the agent can actually run to know it's done — the single highest-leverage line in an agentic prompt.

Why it matters — The single highest-leverage line in a brief. The vendor's own rule for long-running agents: mark work complete 'only after end-to-end verification confirms it works, not when the code is written.' A done-check the agent can run converts 'I finished' into 'it passed.' (Anthropic memory-tool guidance, verbatim)

Option	Pick when	Emits into the prompt
Totals reconcile	data jobs	totals reconcile exactly to <code>{{trusted anchor}}</code>
Tests pass	code	the test suite passes from clean
Files re-open clean	file creation	every deliverable re-opens without errors and matches the spec above
Checklist satisfied	always available	every item in is checked and evidenced
◇ Rendered and inspected	slides, documents, dashboards	outputs are rendered to images and visually inspected against the requirement
🔗 custom	always available	Custom verification...

⚠ Anti-patterns: a mission with no runnable check · 'improve X' with no definition of improved

Context `class AgentContext extends Context // coverage pressure`

Tribal knowledge, written down — the agent steps on every unstated quirk. A chat answer touches your domain once; an agent touches it at every step. The glossary becomes a controlled vocabulary; every quirk carries its rule.

Prevents: silent invented fixes — the most expensive agent error, because they look like diligence.

Domain in three sentences `AgentContext.domain : string [free text]`

The process, the actors, what 'good' means.

Why it matters — Three sentences of process, actors, and what 'good' means give every downstream judgment a frame. Without them the agent optimizes for generic best practice — generically right, specifically wrong for your operation. (context-element mechanics)

Option	Pick when	Emits into the prompt
🔗 custom	always available	What this process is, who's involved, and what good looks like...

Glossary (controlled vocabulary) `AgentContext.glossary : Map [pick any]`

Every metric as a card: numerator, denominator, grain, inclusions, exclusions. The agent will otherwise compute three subtly different versions.

Why it matters — An agent touches every source repeatedly: a guessed denominator doesn't make one wrong sentence — it propagates through every transformation, chart, and slide of the run. Metric cards make definitions load-bearing once and correct everywhere. (source-register discipline; pipeline field practice)

Option	Pick when	Emits into the prompt
Metric card skeleton	every presented metric	<code>{{METRIC}} = {{numerator}} ÷ {{denominator}} at {{grain}}; include {{...}}; exclude {{...}}; zero denominator → NULL, excluded from averages; target {{...}}.</code>
🔗 custom	always available	TERM = definition...

Entities & identifiers `AgentContext.entities : Registry [free text]`

Part numbers, site names, aliases, vendor codes — spelled the way the audience spells them.

Why it matters — Exact identifiers, spelled the way the audience spells them, prevent silent fuzzy matching — the join on 'almost the same part number' that poisons a dataset without raising an error. Enumerated entities turn matching into lookup. (data-pipeline practice)

Option	Pick when	Emits into the prompt
🔗 custom	always available	e.g., <code>'{{id}} = {{meaning}}; {{alias}} → {{canonical}}'</code>

Known quirks → rules `AgentContext.quirks : List [pick any]`

Things that have bitten before, each with the rule that disarms it.

Why it matters — Each known quirk with its disarming rule converts institutional scar tissue into instructions. Undocumented, the same quirk gets 'fixed' differently every run — plausibly, silently, and sometimes destructively. (hallucination mechanics applied to data)

Option	Pick when	Emits into the prompt
Typed totals	manual workbooks	Some sheets carry typed totals rows – detect, exclude, recompute, and report disagreement.
Aliased IDs	multi-system data	{{ID A}} appears as {{ID B}} in {{source}} – map via the reference table, flag unmapped as HARD.
Re-inspection double-count	process data with rework	Re-processed items reappear – count once, at first occurrence, per the carry-over rule.
◇ Grain differs by source	mixed-grain sources	{{Source A}} is unit-level, {{Source B}} is lot-level – never mix grains in one denominator; keep separate layers.
🔧 custom	always available	quirk → rule...

Targets & prior work `AgentContext.targets, AgentContext.priorWork [free text]`

Thresholds that matter, and files or decisions the output must stay consistent with.

Why it matters — Thresholds and previously-published numbers are the run's consistency constraints: they anchor new output to figures the organization already trusts and make drift a detectable event instead of a surprise in a meeting. (reconciliation-anchor family)

Option	Pick when	Emits into the prompt
🔧 custom	always available	targets: ... · stay consistent with: {{file/page}} (definitions, colors, names)

△ Anti-patterns: definitions the agent is left to infer · quirks that live only in someone's head

Environment `class Environment extends PromptElement // NEW IN KIND`

The workspace contract: where the agent may read, where it must write, what is forbidden. Least privilege as prose — and where possible, enforced by sandbox too. The prompt is policy; the sandbox is physics. Prevents: the agent 'helpfully' editing your master file or wandering outside its folder.

Workspace `Environment.workspace : Path [free text]`

Repo/folder and branch. One initiative, one folder.

Why it matters — One initiative, one folder is blast-radius design: everything the run touches is enumerable, backupable, revertable. It also makes sessions reproducible — harness practice starts every session by re-reading the same workspace. (Anthropic harness: per-session ritual)

Option	Pick when	Emits into the prompt
🔧 custom	always available	e.g., '{{owner/repo}} · branch {{claudeslug}} · folder {{initiative}}/'

Read/write boundaries `Environment.boundaries : {readOnly, writeTo, scratch} [pick any]`

The three-zone layout: immutable inputs, versioned outputs, regenerable scratch.

Why it matters — The three-zone layout engineers reversibility in before any action runs: immutable inputs can't be corrupted, versioned outputs can't be silently overwritten, scratch is disposable by contract. Vendor tooling enforces the same idea as protected paths and blocked destructive commands. (Claude Code permission modes; git-revert harness practice)

Option	Pick when	Emits into the prompt
Inputs read-only	always	Inputs live in <code>{{inputs/}}</code> and are READ-ONLY – never modify, overwrite, or 'fix' anything there.
Versioned outputs	always	Write deliverables to <code>{{deliverables/}}</code> with a <code>_MMDDYY</code> suffix – never overwrite an input or a previously delivered file.
Scratch zone	multi-step jobs	Intermediates go to <code>{{work/}}</code> – everything there may be regenerated at any time.
🔧 custom	always available	read-only: <code>{{path}}</code> · write to: <code>{{path}}</code> · scratch: <code>{{path}}</code>

Tools allowed `Environment.tools : Set [pick any]`

Name what exists. Anything unnamed is a guess.

Why it matters — Least privilege starts with naming what exists — anything unnamed is a guess the agent may make for you. Published guidance risk-rates every tool (low/medium/high) by write access, reversibility, and impact; the allowed list is where that rating lives. (OpenAI practical guide, tool safeguards)

Option	Pick when	Emits into the prompt
File system	file jobs	file system (read/write within boundaries)
Shell	automation	shell
Python + data stack	data/deliverable jobs	python 3 with pandas, openpyxl, python-pptx, matplotlib
Web search	research	web search and fetch
◇ Connected apps	integrations in scope	connected apps: <code>{{Drive / Notion / email – name them}}</code>
🔧 custom	always available	Custom tool...

Forbidden actions `Environment.forbidden : Set [pick any]`

The explicit denylist — name what would be assumed available otherwise.

Why it matters — An explicit denylist beats an implied one because evaluation order is deny-first: what is forbidden stays forbidden regardless of what else allows it. Name what would otherwise be assumed available — email, deletion, posting. (permission-rule precedence: deny > ask > allow)

Option	Pick when	Emits into the prompt
No network	confidential local data	make any network calls
No sending	almost always	send email or messages to anyone
No outside writes	always	write anywhere outside the workspace
No deletions	default-safe	delete anything – flag instead
No new dependencies	controlled environments	install tools or dependencies without asking
◇ No live systems	enterprise systems exist	query or touch live production systems (<code>{{SAP / MES / QMS}}</code>)
🔧 custom	always available	Custom prohibition...

Confidentiality `Environment.confidentiality : Policy [pick any]`

What may leave the workspace, and in what form.

Why it matters — The documented exfiltration pattern needs three legs at once: private data, untrusted content, and an outbound channel. The confidentiality clause removes or narrows the outbound leg — say what may leave the workspace, and in what form. (lethal-trifecta analysis; PII-filter guardrails)

Option	Pick when	Emits into the prompt
Summaries only	internal data with a reporting surface	Summaries may be posted to <code>{{destination}}</code> ; raw rows and names never leave the workspace.
Nothing leaves	restricted material	Nothing from this workspace is posted, shared, or transmitted anywhere.
No PII in outputs	people data touched	No personal data appears in any output, including examples.
 custom	always available	Custom rule...

Action risk policy `Environment.riskPolicy : Set // OpenAI tool-risk ratings; Magentic-UI reversibility classes [pick any]`

How actions are triaged by consequence — the published patterns.

Why it matters — Triaging actions by consequence is the published pattern on both major stacks: reversible proceeds, irreversible asks — operationalized as low/medium/high tool ratings and never/maybe/always-irreversible action classes. (OpenAI practical guide; Magentic-UI action guards)

Option	Pick when	Emits into the prompt
Risk-rate tools	mixed tool sets	Treat tools by risk (read-only vs write, reversibility, permissions, financial impact): low-risk freely, medium-risk with logging, high-risk only after asking.
Reversibility triage	always sound	Triage every action: never-irreversible → just do it; maybe-irreversible → assess, and ask if unsure; always-irreversible (delete, send, publish, pay) → explicit approval first.
◇ Sandbox assumed	sandboxed runs	Assume filesystem and network isolation are enforced — if an action fails on a boundary, that is the boundary working; report it, don't work around it.
 custom	always available	Custom risk rule...

sandbox() — Where the tool supports it, enforce boundaries technically (filesystem/network isolation, permission modes) — the prompt states policy, the sandbox enforces it. Use when: any unattended run.

△ Anti-patterns: 'files are in the repo' (no zones) · an unstated tool assumption ('it surely won't email anyone')

Inputs `class Inputs extends PromptElement // NEW IN KIND`

The source register: each source's location, format, grain, trust level, and known issues. The grain line — 'one row = what' — determines every denominator downstream. Most wrong numbers are born as inferred schemas.

Prevents: the agent guessing your schema and silently mixing grains.

Source register `Inputs.sources : Source[] [free text]`

One block per source. The skeleton below is the discipline — fill every line or write UNKNOWN.

Why it matters — A source register makes every input's location, format, grain, and meaning explicit before anything is computed — profiling before processing. Skipped, the first wrong assumption about what one row means propagates through the entire run. (ETL discipline)

Option	Pick when	Emits into the prompt
 custom	always available	S1 — <code>{{name}}</code> · location <code>{{path}}</code> · format <code>{{xlsx sheet 'X', header row N}}</code> · one row = <code>{{grain}}</code> · key <code>{{cols}}</code> · coverage <code>{{dates/sites}}</code> · owner <code>{{who}}</code> · known issues <code>{{...}}</code>

Trust levels `Source.trust : enum [pick any]`

Declare per source — trust level dictates validation depth.

Why it matters — Trust level dictates validation depth — the data-side version of tool risk ratings: a system export and a hand-typed tracker must not get the same scrutiny. Declaring it routes checking effort to where errors actually live. (risk-rating pattern applied to sources)

Option	Pick when	Emits into the prompt
Authoritative	systems of record	{{source}}: authoritative – other sources reconcile TO it.
Derived	reports built from other data	{{source}}: derived – trace to its origin before trusting.
Manual entry	hand-filled sheets	{{source}}: manual entry – re-derive every total from rows; expect typos; reconcile everything.
 custom	always available	e.g., 'S1 authoritative · S2 manual entry'

The UNKNOWN rule `Inputs.unknownPolicy : Rule [pick one]`

The register's cardinal rule.

Why it matters — 'Write UNKNOWN, never guess' turns the model's strongest failure mode — plausible gap-filling — into a logged, owned item. An UNKNOWN is a work item; a guess is a landmine with a delay. (abstention research applied to pipelines)

Option	Pick when	Emits into the prompt
UNKNOWN, never guessed	always on	Any fact about an input you cannot establish: write UNKNOWN and raise it as an open item before Phase 1 ends – never guess.

profile_first() — Before processing: list sheets, headers, row counts, date ranges, null rates, duplicate keys — and report the structure for confirmation. Catches surprises at the cheap end. Use when: every data job.

△ Anti-patterns: a file name standing in for a source description · mixing unit-level and lot-level sources in one metric

Plan `class Plan extends PromptElement // NEW IN KIND`

Bounded method: the steps in order, methods allowed, methods requiring permission. Specific enough to prevent improvised rigor, open enough to adapt to what it finds — the right-altitude element. **Prevents: creative methodology at step 14 of 20, discovered only at review.**

Steps `Plan.steps : OrderedList [free text]`

Each step: what it consumes, what it produces.

Why it matters — Steps with named inputs and outputs make intermediates inspectable — the documented reason explicit pipelines still matter: 'useful when you need to inspect intermediate outputs or enforce a specific pipeline structure.' Audit trails are built here. (Anthropic chaining guidance)

Option	Pick when	Emits into the prompt
 custom	always available	1. {{step}} → {{produces}} 2. {{step}} → {{produces}} ...

Methods allowed `Plan.allowed : Set [pick any]`

Boring and reproducible. The default analytical toolkit.

Why it matters — Boring, reproducible methods are a feature: descriptive stats and run charts re-run identically and survive challenge; clever one-offs don't. Compute-via-code is non-negotiable — model arithmetic degrades sharply as the numbers grow. (Faith & Fate, NeurIPS 2023)

Option	Pick when	Emits into the prompt
Descriptive stats	data	descriptive statistics
Pareto	driver questions	Pareto analysis
Run charts	trends	run charts; control limits only with ≥ 10 points – otherwise say 'too few points for limits'
Rolling averages	noisy series	rolling averages, window stated
Proportions with n	rates	simple proportions, always with n
 custom	always available	Custom method...

Not without asking `Plan.restricted : Set [pick any]`

Propose at a gate; never just use.

Why it matters — 'Propose at a gate, never just use' routes method escalation through a human — the high-risk and failure-threshold triggers vendors place on agent actions, applied to analysis choices that would change what the numbers mean. (OpenAI escalation guidance)

Option	Pick when	Emits into the prompt
No causal claims	observational data	causal claims – 'associated with' only, never 'caused by'
No models/forecasts	unless explicitly in scope	forecasting or model fits of any kind
No data dropping	always	dropping data to clean up a trend
No scope additions	always	adding analyses beyond the numbered questions
 custom	always available	Custom restriction...

△ Anti-patterns: 'do it step by step' (no steps given) · unlimited method freedom on a governed deliverable

Checks `class Checks extends PromptElement // NEW IN KIND – the honesty engine`

Machine-evaluable verification: HARD stops the run, SOFT flags and continues. Write checks as equalities and set-memberships the agent can actually evaluate — and anchor to one number you already trust. Without an anchor, every check can pass on the wrong data. **Prevents: plausible-but-wrong deliverables — the agent's version of hallucination.**

Reconciliation anchor `Checks.anchor : TrustedNumber [free text]`

One number from OUTSIDE the run that the results must tie to. The most important single check.

Why it matters — One number from outside the run that the results must tie to catches whole-pipeline errors — wrong filters, dropped rows, doubled joins — with a single comparison. Recalculate, scan for errors, tie out: the pattern vendor tooling mandates before delivery. (reconciliation/recalc practice)

Option	Pick when	Emits into the prompt
 custom	always available	e.g., 'total units = {{n}} per {{last cycle's master / system report}}'

HARD checks (stop the run) `Checks.hard : Set [pick any]`

Failure means stop, show the offending rows, and wait.

Why it matters — Rules-based feedback is 'the best form of feedback' an agent gets — clear rules, plus which rule failed and why. HARD checks make the run fail loudly at the failure point instead of finishing wrong politely; showing the offending rows makes the failure debuggable. (Anthropic Agent SDK verify-work)

Option	Pick when	Emits into the prompt
Row accounting	any transformation	rows_in = rows_kept + rows_dropped per source, every drop with a logged reason
Reconciliation	always	computed totals match the anchor exactly
Identities	derived metrics	internal identities hold on every row ($\{a\} + \{b\} = \{total\}$; $0 \leq rate \leq 1$)
Key integrity	joined data	no duplicate primary keys; every child key exists in its parent
Domains	categorical data	every coded value within its allowed set after normalization; no UNKNOWN in presented numbers
Files open clean	file deliverables	every deliverable re-opens without errors (<code>{{pptx via python-pptx,.xlsx via openpyxl}}</code>)
◇ Citations resolve	research jobs	every citation/link resolves to a real source that says what's claimed
◇ Recalc & error scan	spreadsheet deliverables	recalculate all formulas and scan for spreadsheet error values (REF, DIV/0, VALUE, N/A) before delivery – zero errors allowed
◇ Smoke test first	long-running / resumed jobs	run the basic end-to-end smoke test BEFORE starting new work – never build on a broken state
◇ End-to-end as a user	apps, dashboards, documents	verify the deliverable end-to-end the way a user would experience it (open it, click it, run it) before marking done
🔧 custom	always available	Custom HARD check...

SOFT checks (flag and continue) `Checks.soft : Set [pick any]`

Suspicion, not failure — logged to open items.

Why it matters — SOFT separates suspicion from failure so the run neither ignores anomalies nor cries wolf: flagged items land in open items with evidence attached, and a human triages them at review instead of mid-run. (check-type taxonomy; open-items discipline)

Option	Pick when	Emits into the prompt
Null-rate threshold	recurring data	null rate of $\{key\ field\} \leq \{x\}\%$ – flag any spike vs. last cycle
Outliers listed	measurements	values beyond <code>{{3σ or physical limits}}</code> listed with context – never removed
Small-n tags	proportions	any presented group with $n < \{20\}$ tagged small-n, with n shown on the chart
Delta vs prior	recurring reports	any metric moving $> \{x\}$ vs. the prior cycle gets a one-line explanation or a flag
🔧 custom	always available	Custom SOFT check...

`validation_report()` — All checks print to one reconciliation table (pass/fail per check, soft flags listed) — the artifact a human reviews at the gate. Use when: always.

`check_types()` — The canonical verification trio: rules-based feedback (tests, linting, recalculation — strongest), visual feedback (screenshots/renders), LLM-as-judge (fuzzy criteria — weakest, use last). Use when: choosing what a check should be made of.

⚠️ Anti-patterns: 'check the data quality' (not evaluable) · ten checks, no anchor

Outputs `class Outputs extends Format // + audit trail`

The deliverable contract and the audit trail — exact names, exact locations, three logs. A colleague should be able to find, understand, and regenerate everything from the folder alone, six weeks later. **Prevents: results you can't locate and conclusions you can't trace.**

Deliverables `Outputs.deliverables : File[]` [free text]

Numbered, with the `_MMDDYY` convention and what each contains.

Why it matters — Numbered deliverables with exact names make 'done' countable — the reason harness guidance tracks work as a structured list with pass flags rather than prose: it prevents premature completion and orphaned extras alike. (Anthropic harness: feature-list pattern)

Option	Pick when	Emits into the prompt
 custom	always available	1. deliverables/{{NAME}}_MMDDYY.{{ext}} - {{contents}} 2. ...

Logs `Outputs.logs : Set` [pick any]

The three files that make a run auditable.

Why it matters — CHANGES, FINDINGS, and OPEN_ITEMS make a run auditable after the fact — the 'structured updates' long-running-agent guidance requires every session to leave. Without logs there is no way to distinguish what the agent did from what it says it did. (Anthropic harness guidance)

Option	Pick when	Emits into the prompt
CHANGES.md	always	append to CHANGES.md: date, what changed, why, who asked
FINDINGS.md	analysis	write FINDINGS.md: each claim → evidence (table/figure ref) → confidence (high/med/low and why)
OPEN_ITEMS.md	always	maintain OPEN_ITEMS.md: assumptions, exclusions, questions — each with an owner and status
 custom	always available	Custom log...

Reproducibility `Outputs.repro : Requirement` [pick any]

The regeneration guarantee.

Why it matters — The regeneration guarantee — same inputs, same command, same outputs — turns a deliverable into evidence. It is why outputs carry formulas rather than pasted values and why environments get an init script: reproducibility is engineered, not hoped for. (harness init-script pattern; workbook discipline)

Option	Pick when	Emits into the prompt
One-command rebuild	data pipelines	a single script regenerates every deliverable from the raw copies with no manual steps
Numbered scripts	multi-phase jobs	scripts numbered by phase (01_extract ... 05_present) in <code>{{work/scripts/}}</code>
 Numbers trace	governed reporting	<code>numbers_trace.csv</code> maps every presented number to its source cell
 custom	always available	Custom requirement...

⚠ Anti-patterns: 'save the results' (where? named what?) · no way to regenerate after a source correction

Process & Gates `class Process extends TheStop` // matured: phases + gates + autonomy

Phases with human checkpoints at the cheap-error points — and the autonomy rules between them. Gate 1 catches definition errors (a sentence to fix), Gate 2 data errors, Gate 3 message errors (a retraction to fix). Without autonomy rules the agent asks about everything or nothing. **Prevents: compounding — hours of competent work on a wrong foundation.**

Phases `Process.phases : OrderedList` [free text]

Phase 0 is always intake: the agent restates the mission and its assumptions in its own words.

Why it matters — Phase 0 — the agent restates the mission and its assumptions in its own words — is the cheapest error catch in the run: a misread brief costs pennies here and a full rerun later. Plan-then-approve is the documented pattern across vendor tooling. (Claude Code plan mode; Magentic-UI co-planning)

Option	Pick when	Emits into the prompt
 custom	always available	Phase 0 intake → Phase 1 <code>{{...}}</code> → Phase 2 <code>{{...}}</code> → final QA phase

Gates `Process.gates : Set [pick any]`

Human checkpoints, in ascending cost of the error they catch.

Why it matters — Gates put human judgment where published guidance puts it: before edits (plan approval), before irreversible or high-stakes actions, at failure thresholds, and at final review — each gate catching its error class at its cheapest point. (OpenAI practical guide; gate-placement survey)

Option	Pick when	Emits into the prompt
Gate 1 — Definitions	every data/document job	► GATE 1 (definitions): restate mission, metrics, schema, and assumptions; ask only what blocks Phase 1. Do not touch data before my confirmation.
Gate 2 — Reconciliation	data jobs	► GATE 2 (reconciliation): present the validation table. Any HARD fail stops here regardless.
Gate 3 — Headlines	deliverables people will see	► GATE 3 (headlines): titles and key numbers as plain text for approval BEFORE rendering anything.
Gate — Pre-send	anything outward-facing	► GATE (pre-send): show exactly what will be sent/published and to whom; wait for explicit approval.
◇ Gate — Sample first	bulk pipelines	► GATE (sample): process <code>{{3}}</code> representative items and show the result before batch-processing the rest.
🔧 custom	always available	Custom gate...

Autonomy rules `Process.autonomy : {proceed, stopIf[]} [pick any]`

Both halves matter: without 'proceed' you built a slow chatbot; without 'stop' you built a liability.

Why it matters — Both halves are load-bearing: without proceed rules you built a slow chatbot that asks about everything; without stop rules you built a liability that asks about nothing. The published stop triggers are high-risk actions and exceeded failure thresholds. (OpenAI human-intervention guidance)

Option	Pick when	Emits into the prompt
Proceed between gates	always include	Between gates, proceed without asking.
Stop: HARD fail	always include	Stop immediately if a HARD check fails – report with the offending rows.
Stop: input mismatch	always include	Stop if an input is missing or shaped differently than described in .
Stop: ambiguity	always include	Stop if a metric definition or instruction is ambiguous – ask, don't interpret.
Stop: irreversible	always include	Stop before anything hard to reverse: deleting, sending, publishing, overwriting. Reversible → act; irreversible → ask.
Stop: published number	recurring reporting	Stop if a result would change a previously published figure.
◇ Retry policy	network/flaky tooling	On transient errors: retry twice with backoff, then stop and report – never silently skip.
🔧 custom	always available	Custom rule...

Autonomy level `Process.autonomyLevel : enum // Feng et al. 2025 – named by the role YOU take [pick one]`
A deliberate design choice, separate from capability. Named by the role you keep — pick the lowest level that still gets the job done.

Why it matters — Autonomy is 'a deliberate design decision, separate from capability' — the published levels are named by the role the human keeps: operator, collaborator, consultant, approver, observer. Pick the lowest level that gets the job done; capability never obliges autonomy. (Feng et al. 2025)

Option	Pick when	Emits into the prompt
Operator (L1)	new job types, low trust so far	Autonomy: OPERATOR level – I approve actions before you execute them; you assist on the sub-tasks I invoke.
Collaborator (L2)	established jobs, active session	Autonomy: COLLABORATOR level – we plan jointly; you execute assigned tasks independently and report progress and blockers transparently.
Consultant (L3)	trusted recurring work	Autonomy: CONSULTANT level – you plan and execute; consult me proactively for preferences and domain judgment.
Approver (L4)	mature briefs with strong checks	Autonomy: APPROVER level – work autonomously on low-stakes steps; bring me only consequential actions and blockers.
◇ Observer (L5)	only with sandboxing + proven track record	Autonomy: OBSERVER level – full autonomy; I monitor the logs. (Requires airtight checks and a contained environment.)

△ Anti-patterns: gates without autonomy rules (ask-everything or ask-nothing) · no Phase-0 restatement (drift caught late)

Rules `class Rules extends PromptElement // invariants`

The non-negotiables — short, absolute, few — that must survive a long context. Rules hold everywhere; checks are evaluated at points. Keep rules under ten lines or they dilute each other. Prevents: standards drift in hour three of a run.

Invariants `Rules.invariants : Set [pick any]`

The standard set — prune to what this job actually needs.

Why it matters — Invariants are the deterministic guardrail class — rules enforced as rules, not judgment calls: never invent data, never overwrite inputs, version everything. Deterministic protections are their own category in vendor guardrail taxonomies precisely because they don't degrade under pressure. (OpenAI rules-based protections)

Option	Pick when	Emits into the prompt
Never invent	always	Never invent, estimate, or back-fill a number. Missing is missing; say so where it matters.
Immutable sources	always	Source files are immutable. Outputs are versioned. Everything regenerable lives in scratch.
Logged transformations	data	Every transformation is a logged rule with a row count. No silent fixes, no manual cell edits.
Assumptions declared	always	State assumptions as assumptions, in OPEN_ITEMS.md – and in the output where they affect a conclusion.
Boring methods	always	Prefer boring, reproducible methods. Propose fancier ones at a gate; don't just use them.
House conventions	org standards exist	Conventions: dates {{ISO in files, MM/DD/YY in prose}}; units {{...}}; language {{...}}.
🔗 custom	always available	Custom rule...

△ Anti-patterns: twenty rules (they dilute) · rules that duplicate checks (keep them distinct)

Quality Bar `class QualityBar extends PromptElement // done, auditable`

The definition of done as a checklist the agent must satisfy — and you can audit in sixty seconds. Checks verify the work mid-run; the quality bar verifies completeness at the end. Distinct jobs. Prevents: 'done' meaning 'I stopped.'

Definition of done `QualityBar.items : Checklist [pick any]`

Tick-able, evidence-able items.

Why it matters — A definition of done as tick-able, evidence-able items is the anti-premature-completion device: 'passes' flips only on verification, and the checks themselves are out of bounds to edit — 'it is unacceptable to remove or edit tests.' (Anthropic harness, verbatim)

Option	Pick when	Emits into the prompt
Placeholders resolved	always	[] Every {{placeholder}} in this brief was resolved or logged as an open item.
Checks pass	always	[] All HARD checks pass; every SOFT flag has a note.
Deliverables verified	file outputs	[] Every deliverable re-opened and verified after creation – not assumed.
Logs written	always	[] CHANGES / FINDINGS / OPEN_ITEMS written; every open item has an owner and status.
Traceability	governed reporting	[] Every presented number traces to a source cell (numbers_trace).
◇ Charts honest	chart deliverables	[] Every chart: title states the finding, axes have units, n shown on proportions, target line where a target exists, source footer.
🔗 custom	always available	Custom item...

△ Anti-patterns: a quality bar the agent can satisfy without evidence

Reporting `class Reporting extends Format // the interface back to you`

The fixed shape of the final message: status, headline, evidence, open items — under 25 lines. Evidence, not assertions: the test it ran and what it returned, not 'everything went well.' Details live in the files, not the message.

Prevents: a narrative essay where a status should be.

Report shape `Reporting.shape : OrderedList [pick any]`

The standard six lines — deselect only what genuinely doesn't apply.

Why it matters — A fixed report shape is why status reads in 30 seconds: the same lines in the same order every run, skimmable by humans and parseable by tools. Every agent SDK ships fixed result-status enums for the same reason — variable-shape reports hide failures. (SDK result/status enums)

Option	Pick when	Emits into the prompt
1 • Status	always	1. Status: DONE STOPPED AT GATE n BLOCKED – one line why.
2 • Headline	always	2. Headline answer with the number(s) – one sentence.
3 • Deliverables	always	3. Deliverables – paths.
4 • Checks	always	4. Checks summary – what passed, what was flagged.
5 • Open items	always	5. Open items needing a human – who, what, by when.
6 • Risks	always	6. What I'd check next / risks to the conclusion – at most 3 bullets.
🔗 custom	always available	Custom section...

Report discipline `Reporting.rules : Set [pick any]`

The rules that keep reports readable and honest.

Why it matters — Evidence over assertion: 'done' means the check it ran, the output it got, and the file re-opened and verified — not the claim. Reports that cite their evidence convert trust from personality to track record, which is what agent oversight is supposed to run on. (harness verification guidance; NCSC oversight principle)

Option	Pick when	Emits into the prompt
25-line cap	always	Keep the report under 25 lines. Details live in the files, not the message.
Evidence required	always	Claims of success carry evidence: the command run and its output, or the re-opened file.

△ Anti-patterns: a report that narrates the work instead of answering the question

PART 3 — INSTANCES

The template library as saved configurations

The PRIMARY path into the tools is the five course task families (Analyze data · Build a dashboard · Prepare a presentation · Summarize email · Explain clearly) — no template code needed. Each legacy library template below remains available as an optional instance of the class system — a saved set of element choices. Load any of these in the Template Creator (the "Start from a template" menu) and inspect which options it picked; that inspection is the fastest way to learn the taxonomy.

Instance	Mode	Elements in play	Options pre-selected
G1 · Writing — draft from raw material	generative	5	15 options + 2 custom blocks
G2 · Data analysis — trustworthy numbers	generative	6	18 options + 3 custom blocks
G3 · Evaluation — rubric grading	generative	5	9 options + 2 custom blocks
G3b · Blind critique — anti-sycophancy	generative	4	5 options + 2 custom blocks
G4 · Spreadsheet — build a workbook	generative	5	9 options + 2 custom blocks
G5 · Presentation — outline first	generative	5	11 options + 2 custom blocks
G6 · Interactive HTML — dashboard / calculator	generative	5	4 options + 4 custom blocks
G7 · Research brief — cited	generative	5	9 options + 3 custom blocks
G8 · Documents — summarize / compare	generative	5	8 options + 2 custom blocks
A1 · Mission brief — full skeleton	agentic	12	46 options + 10 custom blocks
A2 · ETL → Presentation — data to defensible slide	agentic	12	75 options + 8 custom blocks
A3 · Recurring cycle — reuse by diff	agentic	5	11 options + 4 custom blocks
A4 · Workspace standing rules (CLAUDE.md)	agentic	3	11 options + 4 custom blocks
A5 · Document pipeline — folder → tracker	agentic	8	20 options + 6 custom blocks
GX · Comparison study → decision deck	generative	5	5 options + 2 custom blocks
GX · Scored comparison workbook	generative	5	6 options + 1 custom blocks
AX · Versioned model/report update	agentic	5	7 options + 4 custom blocks
GX · Master document → audience variants	generative	5	5 options + 2 custom blocks

APPENDIX — THE DECOMPOSITION, VALIDATED AGAINST CODE

Engineering systems already model prompts this way

Our concept	DSPy	LangChain	Vendor APIs
Element (class)	Signature class (docstring = task instructions)	ChatPromptTemplate (message roles)	system / developer / user message roles
Attribute (property)	InputField / OutputField — “field naming is the cheapest optimization”	template variables	request parameters (tone, format live in the text)
Option (allowed value)	Literal["spring","summer",...] — typed enums the LM must respect	few-shot example selectors	structured outputs / JSON schema enums
Technique (method)	Module: Predict, ChainOfThought, ReAct — strategy separate from contract	chains / runnables	reasoning-effort & verbosity parameters
Template (instance)	a compiled program	a filled PromptTemplate	a saved prompt / GPT / Gem / skill

The separation this taxonomy makes — WHAT you want (elements/attributes) vs. HOW the model should work on it (techniques) — is the same separation DSPy makes between Signatures and Modules, and the pattern literature makes between a pattern’s “fundamental contextual statements” and its usage. The course teaches the HOW as one loop (PDCA) rather than a numbered technique list: every technique in this reference is a method inside Plan,

Do, Check, or Act. The full sources live in notes/research/ (t1-t6).

ELEMENTS AS REQUIREMENT TYPES

Every element in this taxonomy specifies requirements for a reply or artifact. The mapping below connects the anatomy to the requirement-type menu used across the course (full per-artifact menus with instruction and acceptance examples: references/Requirements_by_Artifact.md and the Field Guide, part 6):

You write...	You are specifying...
Role	role & working behavior — a perspective plus observable conduct; a title alone is weak
Task	functional behavior + audience & use — what it must do, for whom, toward which decision
Context & sources	data & provenance + scope & exclusions — allowed inputs, definitions, period; consequential facts become explicit instructions
Tone	tone & presentation — voice, reading level, units, formats
Format	structure & interface + delivery & compatibility — sections, tabs, controls; file type, editability, offline use
Methods	method & business rules — formulas, denominators, missing-data treatment, precedence
Interaction (agentic)	functional behavior + constraints & authority — controls and behaviors; when to proceed, when to stop and ask
Boundaries (Out & Stop)	content & completeness + constraints & authority — the missing-information rule and the scope edge
Checks	quality & acceptance — observable evidence: reconcile, recompute, exercise, cover